

February 17, 2026

## Task 04: Predicting Insurance Claim Amounts.

Objective: Estimate the medical insurance claim amount based on personal data.

Dataset: Medical Cost Personal Dataset. About Dataset: Context: Machine Learning with R by Brett Lantz is a book that provides an introduction to machine learning using R. As far as I can tell, Packt Publishing does not make its datasets available online unless you buy the book and create a user account which can be a problem if you are checking the book out from the library or borrowing the book from a friend. All of these datasets are in the public domain but simply needed some cleaning up and recoding to match the format in the book. Content: Columns: age: Age of primary beneficiary. sex: Insurance contractor gender, female, male bmi: Body mass index, providing an understanding of body weights that are relatively high or low relative to height, objective index of body weight ( $\text{kg} / \text{m}^2$ ) using the ratio of height to weight, ideally 18.5 to 24.9. children: Number of children covered by health insurance / Number of dependents. smoker: Smoking. region: The beneficiary's residential area in the US, northeast, southeast, southwest, northwest. charges: Individual medical costs billed by health.

Load The Dataset:

[1]: `import pandas as pd`[2]: `Dataset = pd.read_csv("insurance.csv")`

Data Exploration:

[3]: `Dataset.head()`

```
[3]:   age     sex     bmi  children smoker      region    charges
  0   19  female  27.900        0    yes  southwest  16884.92400
  1   18    male  33.770        1    no   southeast  1725.55230
  2   28    male  33.000        3    no   southeast  4449.46200
  3   33    male  22.705        0    no  northwest  21984.47061
  4   32    male  28.880        0    no  northwest  3866.85520
```

[4]: `Dataset.tail()`

```
[4]:      age     sex     bmi  children smoker      region    charges
1333    50    male  30.97        3    no  northwest  10600.5483
1334    18  female  31.92        0    no  northeast  2205.9808
1335    18  female  36.85        0    no  southeast  1629.8335
```

```
1336    21  female  25.80         0     no southwest  2007.9450
1337    61  female  29.07         0    yes northwest  29141.3603
```

[5]: Dataset.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   age         1338 non-null    int64  
 1   sex          1338 non-null    object  
 2   bmi          1338 non-null    float64 
 3   children    1338 non-null    int64  
 4   smoker       1338 non-null    object  
 5   region       1338 non-null    object  
 6   charges      1338 non-null    float64 
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

[6]: Dataset.describe()

```
age            bmi        children      charges
count  1338.000000  1338.000000  1338.000000  1338.000000
mean    39.207025   30.663397   1.094918    13270.422265
std     14.049960   6.098187   1.205493    12110.011237
min     18.000000   15.960000   0.000000    1121.873900
25%    27.000000   26.296250   0.000000    4740.287150
50%    39.000000   30.400000   1.000000    9382.033000
75%    51.000000   34.693750   2.000000    16639.912515
max     64.000000   53.130000   5.000000    63770.428010
```

[7]: Dataset.shape

[7]: (1338, 7)

[8]: Dataset.columns

```
[8]: Index(['age', 'sex', 'bmi', 'children', 'smoker', 'region', 'charges'],
      dtype='object')
```

[9]: Dataset.dtypes

```
[9]: age           int64
      sex          object
      bmi          float64
      children    int64
```

```
smoker      object
region      object
charges     float64
dtype: object
```

Instructions: Visualize how BMI, age and smoking status impact insurance charges.

```
[10]: import matplotlib as pyplot
from matplotlib import pyplot as plt

import seaborn as sns
```

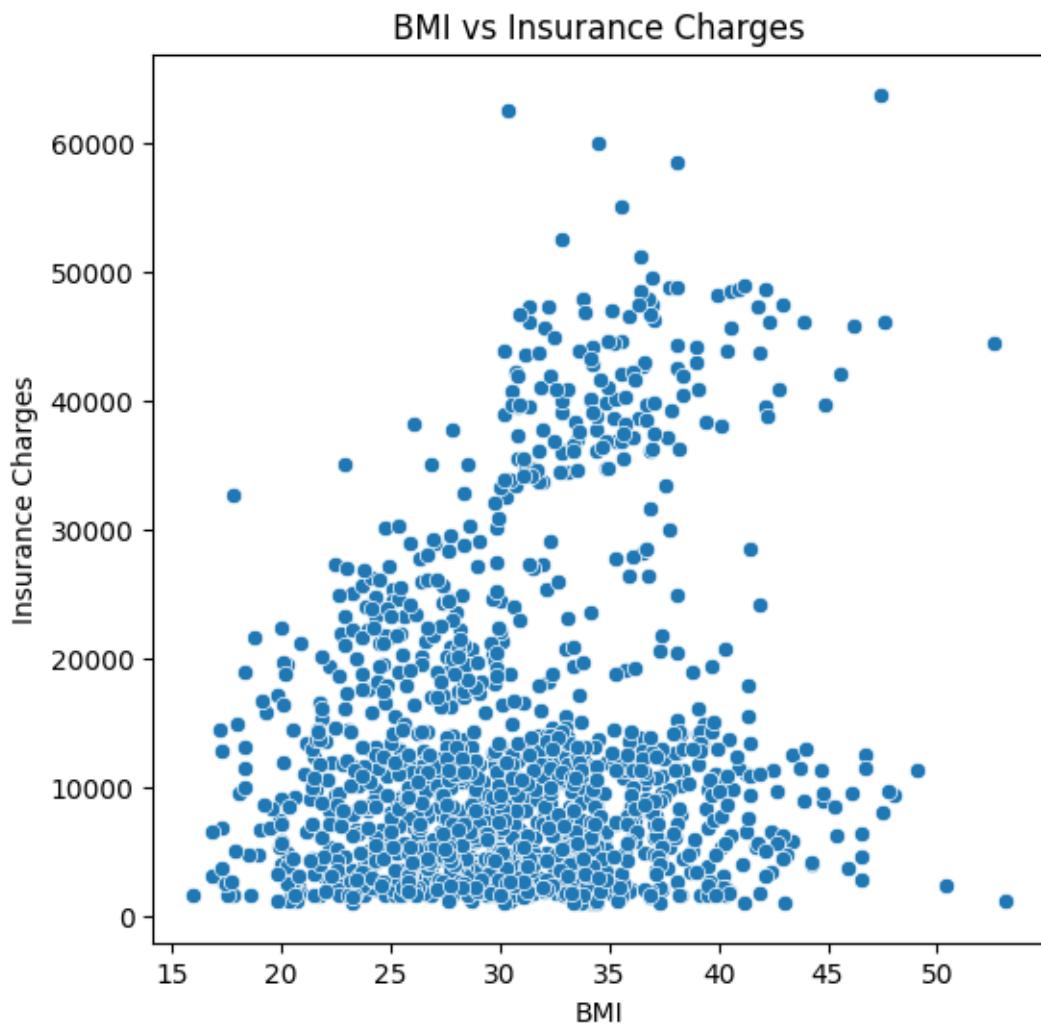
BMI vs Insurance Charges:

```
[11]: plt.figure(figsize=(6,6))

sns.scatterplot(x=Dataset["bmi"] , y=Dataset["charges"])

plt.title("BMI vs Insurance Charges")
plt.xlabel("BMI")
plt.ylabel("Insurance Charges")

plt.show()
```



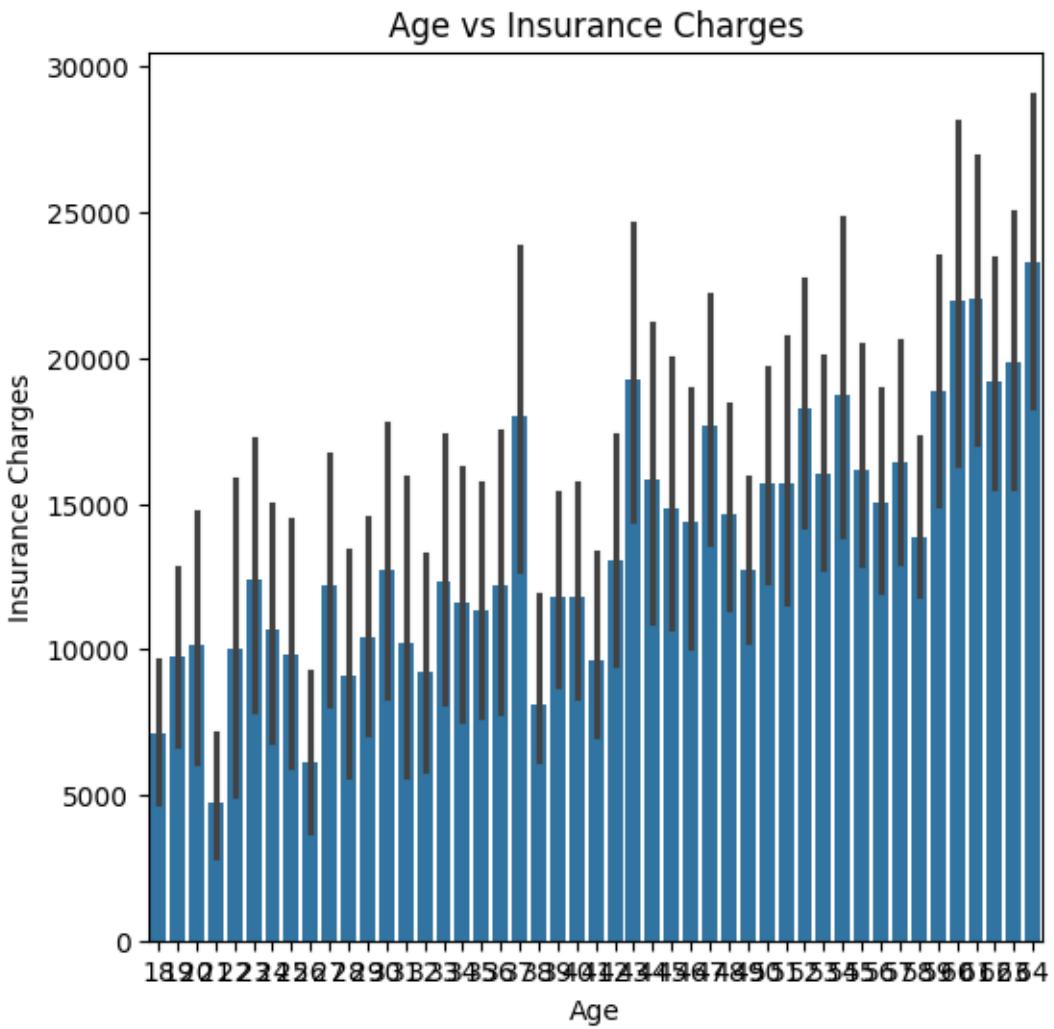
Age vs Insurance Charges:

```
[12]: plt.figure(figsize=(6,6))

sns.barplot(x=Dataset["age"], y=Dataset["charges"])

plt.title("Age vs Insurance Charges")
plt.xlabel("Age")
plt.ylabel("Insurance Charges")

plt.show()
```



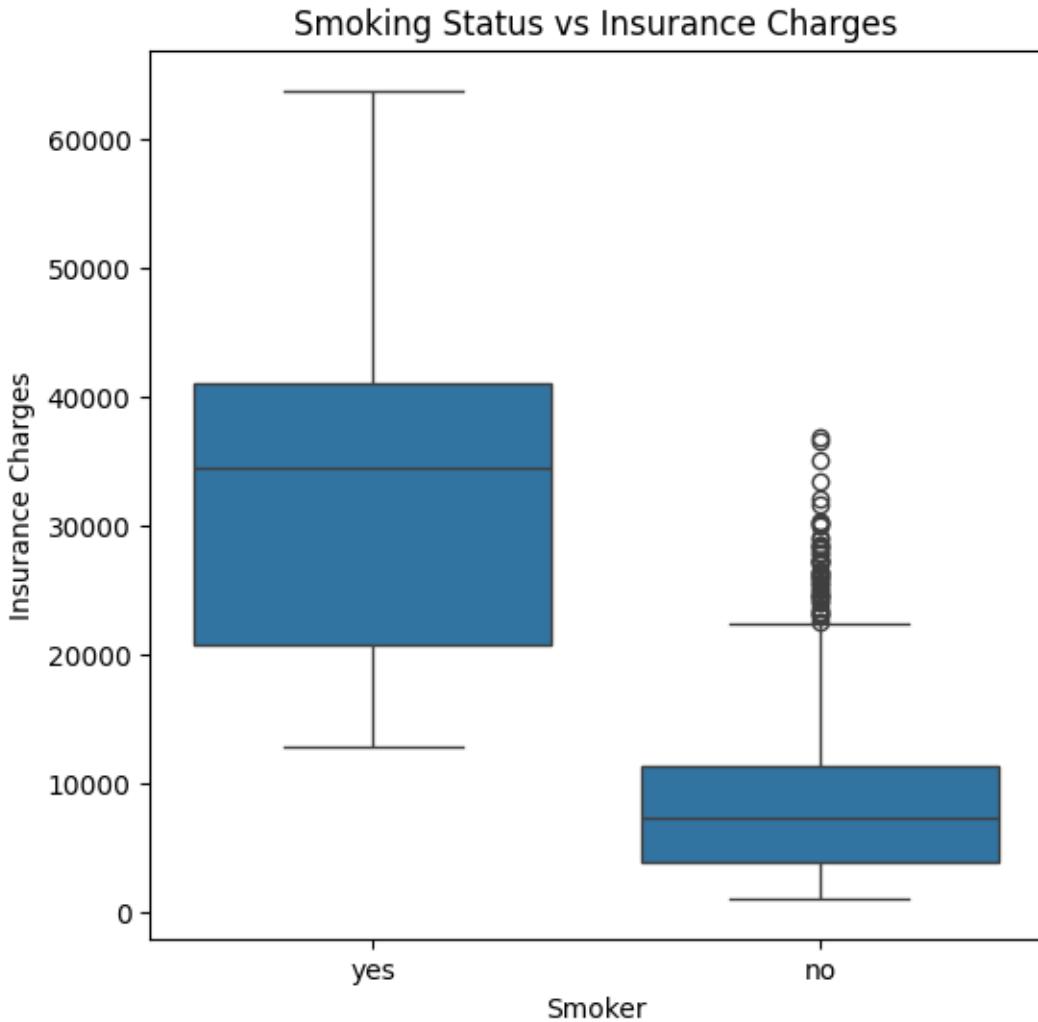
Smoking Status vs Insurance Charges:

```
[13]: plt.figure(figsize=(6,6))

sns.boxplot(x=Dataset["smoker"], y=Dataset["charges"])

plt.title("Smoking Status vs Insurance Charges")
plt.xlabel("Smoker")
plt.ylabel("Insurance Charges")

plt.show()
```



PaitPlot For All The Combinations:

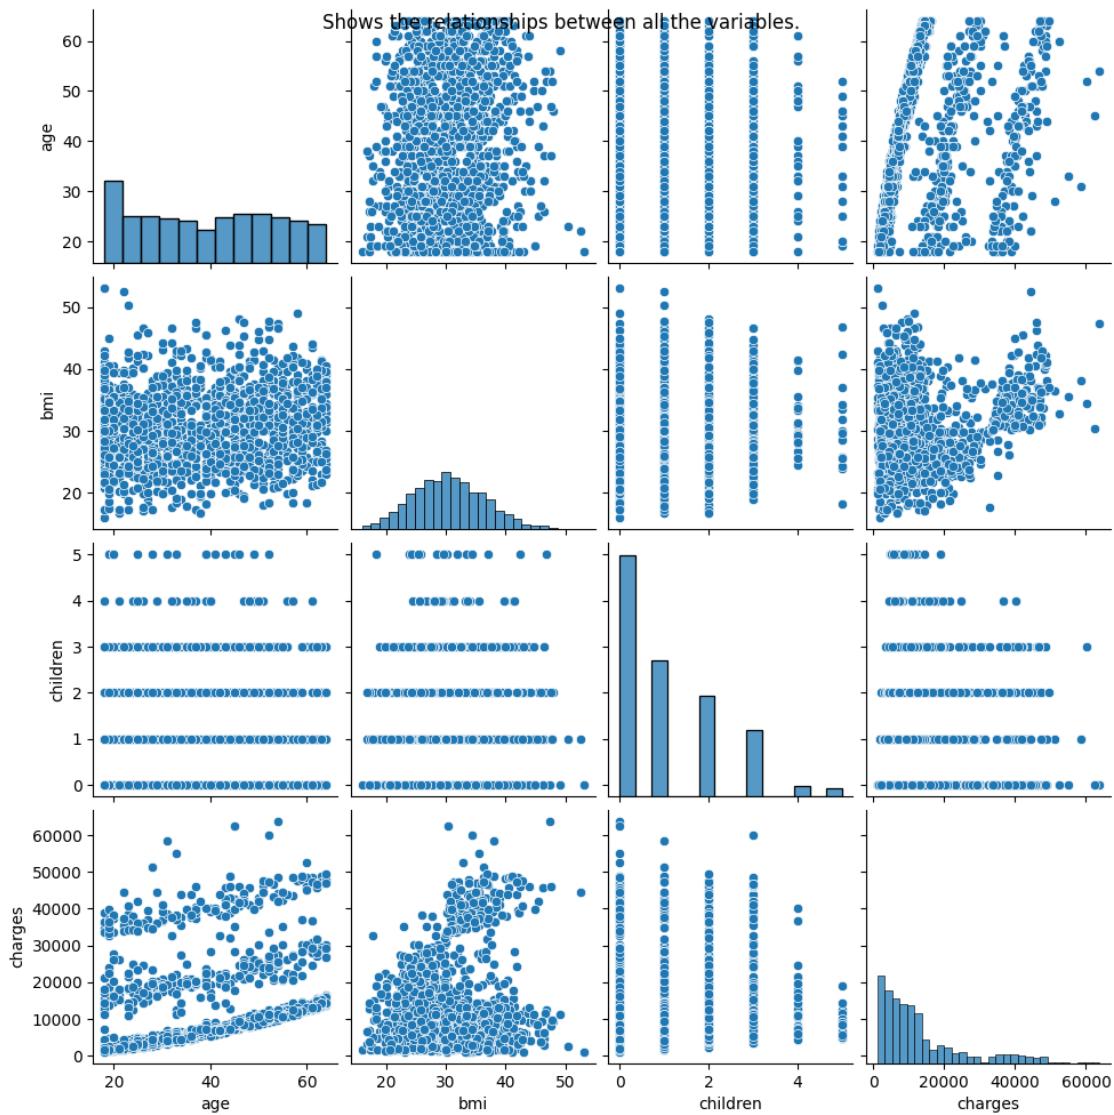
```
[14]: plt.figure(figsize=(6,6))

sns.pairplot(Dataset)

plt.suptitle("Shows the relationships between all the variables.")

plt.show()
```

<Figure size 600x600 with 0 Axes>



Handle Missing values:

```
[15]: Dataset.isnull().sum()
```

```
[15]: age      0
       sex     0
       bmi     0
       children 0
       smoker   0
       region   0
       charges  0
dtype: int64
```

```
[16]: Dataset = pd.get_dummies(Dataset, drop_first=True)
```

```
[17]: Dataset.head()
```

```
[17]:   age      bmi  children      charges  sex_male  smoker_yes  region_northwest \
0    19  27.900         0  16884.92400    False      True        False
1    18  33.770         1   1725.55230     True     False        False
2    28  33.000         3   4449.46200     True     False        False
3    33  22.705         0   21984.47061     True     False        True
4    32  28.880         0   3866.85520     True     False        True

      region_southeast  region_southwest
0                  False                 True
1                  True                False
2                  True                False
3                 False                False
4                 False                False
```

Train a Linear Regression model to predict charges.

```
[18]: import numpy as np

from sklearn.model_selection import train_test_split

from sklearn.linear_model import LinearRegression

from sklearn.metrics import mean_squared_error

from sklearn.metrics import mean_squared_error
```

```
[19]: x = Dataset.drop("charges", axis=1)
y = Dataset["charges"]
```

```
[20]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2,random_state=42)
```

```
[21]: print("Training Data:", x_train.shape)
print("Testing Data:", x_test.shape)
```

Training Data: (1070, 8)  
Testing Data: (268, 8)

```
[22]: Linear_model = LinearRegression()
Linear_model.fit(x_train, y_train)

y_predict_Linear = Linear_model.predict(x_test)
```

```
[23]: model = LinearRegression()
model.fit(x_train, y_train)
```

```
[23]: LinearRegression()
```

Evaluate the Model using the MAE RMSE.

Mean Square Error.

```
[24]: print("Mean Squared Error:", mean_squared_error(y_test, y_predict_Linear))
```

Mean Squared Error: 33596915.851361446

Root Mean Square Error.

```
[25]: rmse = np.sqrt(mean_squared_error(y_test, y_predict_Linear))

print("Root Mean Square Error:", rmse)
```

Root Mean Square Error: 5796.284659276272

```
[26]: comparison = pd.DataFrame({
    "Actual Charges": y_test.values,
    "Predicted Charges": y_predict_Linear})
```

```
[27]: comparison.head()
```

```
[27]:   Actual Charges  Predicted Charges
 0      9095.06825      8969.550274
 1      5272.17580      7068.747443
 2      29330.98315     36858.410912
 3      9301.89355      9454.678501
 4      33750.29180     26973.173457
```

Skills: Regression modeling. Feature correlation and visualization. Error model performance using MAE and RMSE.

Task Completed. Best Wishes. Zaigham Abbas.